

Hands-on 9

CIE5141 Computer Vision in Construction

b07501113 林庭瑄

A. Image Classification

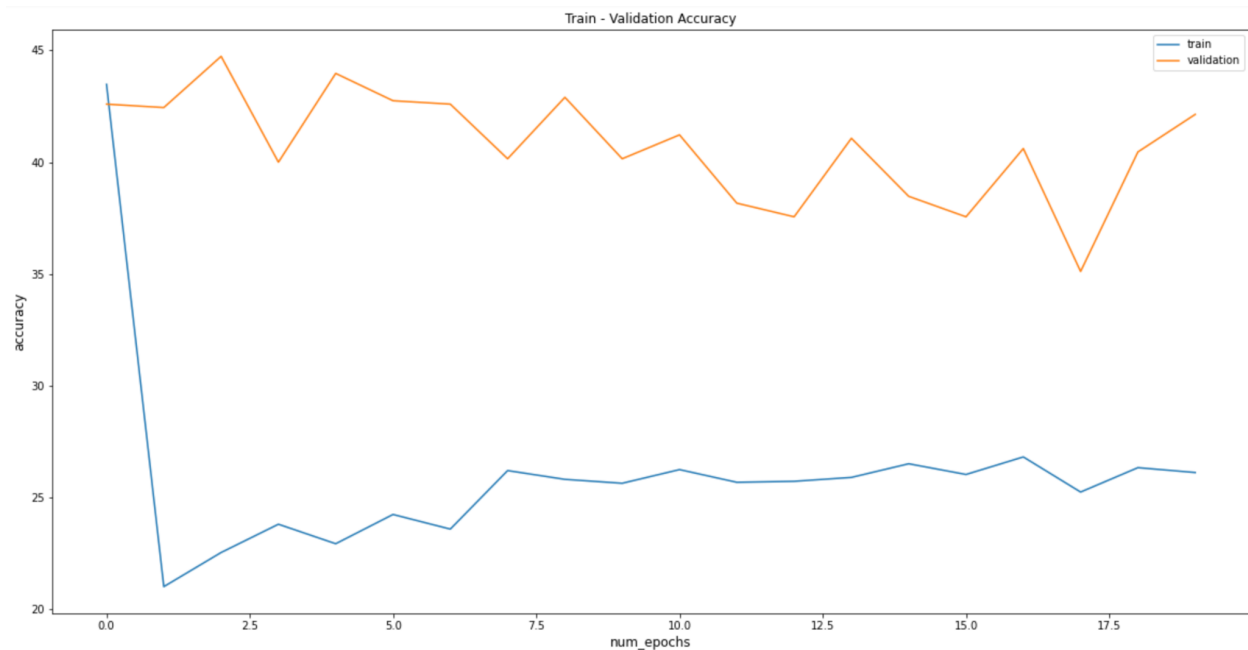
1. Report your results (50%)


(i) Report your accuracy

```
correct=0
total=0
for batch_idx, (data_, target_) in enumerate(test_loader):
    data_, target_ = data_.to(device), target_.to(device) # on GPU
    outputs = model(data_)
    _,pred = torch.max(outputs, dim=1)
    correct += torch.sum(pred==target_).item()
    total += target_.size(0)
print("Accuracy:",correct/total*100)
```


Accuracy: 45.25993883792049

(ii) Paste here the accuracy and loss curves





(iii) Paste here some prediction results



(iv) Mentions the steps taken to improve the accuracy

Data:

70% train
20% validate
10% test

Data normalization & Data augmentation.

Normalize : train&validation&test

RandomCrop: train

Gaussian Blur: train

Random Rotation: train

```
train_transform = transforms.Compose([
    transforms.RandomCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(30, resample=Image.BICUBIC,
expand=False),
    transforms.GaussianBlur(7,3),
    transforms.ToTensor(),
    transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224,
0.225)))
val_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.CenterCrop(224),
    transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224,
0.225)))
test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.CenterCrop(224),
    transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224,
0.225)))
```

Deeper network & Normalization layers

Add several layers

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        # 3 input image channel, 16 output channels, 3x3 square
convolution kernel
        self.conv1 = nn.Conv2d(3,16,kernel_size=3,stroke=2,padding=1)
        self.conv2 = nn.Conv2d(16, 32,kernel_size=3,stroke=2, padding=1)
        self.conv3 = nn.Conv2d(32, 64,kernel_size=3,stroke=2, padding=1)
```

```

        self.conv4 = nn.Conv2d(64, 128, kernel_size=3, stride=2,
padding=1)
        self.conv5 = nn.Conv2d(128, 256, kernel_size=3, stride=2,
padding=1)
        self.conv6 = nn.Conv2d(256, 512, kernel_size=3, stride=2,
padding=1)
        self.conv7 = nn.Conv2d(512, 512, kernel_size=3, stride=2,
padding=1)

        self.pool = nn.MaxPool2d(2, 2)
        self.dropout = nn.Dropout2d(0.4)
        self.batchnorm1 = nn.BatchNorm2d(16)
        self.batchnorm2 = nn.BatchNorm2d(32)
        self.batchnorm3 = nn.BatchNorm2d(64)
        self.batchnorm4 = nn.BatchNorm2d(128)
        self.batchnorm5 = nn.BatchNorm2d(256)
        self.batchnorm6 = nn.BatchNorm2d(512)
        self.fc1 = nn.Linear(64*8, 512 )
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, 128)
        self.fc4 = nn.Linear(128, 64)
        self.fc5 = nn.Linear(64, 32)
        self.fc6 = nn.Linear(32, 13)

def forward(self, x):
    x = self.batchnorm1(F.relu(self.conv1(x)))
    x = self.batchnorm2(F.relu(self.conv2(x)))
    x = self.batchnorm3(self.pool(F.relu(self.conv3(x))))
    x = self.dropout(self.batchnorm3(self.pool(x)))
    x = self.batchnorm4(F.relu(self.conv4(x)))
    x = self.batchnorm5(F.relu(self.conv5(x)))
    x = self.batchnorm6(F.relu(self.conv6(x)))
    x = self.dropout(self.conv7(x))
    x = x.view(-1, 64*8) # Flatten layer
    x = self.dropout(self.fc1(x))
    x = self.dropout(self.fc2(x))
    x = self.dropout(self.fc3(x))
    x = self.dropout(self.fc4(x))
    x = self.dropout(self.fc5(x))
    x = F.log_softmax(self.fc6(x), dim = 1)
    return x

```

Early stopping.

set the epoch to 20 as the validation accuracy starts to decrease.

